
Lab exercise - Part I & II

METHODS AND MODELS FOR COMBINATORIAL OPTIMIZATION



FURNO FRANCESCO

francesco.furno@studenti.unipd.it - 2139507

A.A. 2024 - 2025

Table of contents

| | |
|---|-----------|
| 1. Introduction | 3 |
| 1.1. Problem description | 3 |
| 1.2. Project structure | 3 |
| 1.3. Compilation | 4 |
| 2. Instances generator | 6 |
| 3. Part I: Cplex Solver | 8 |
| 3.1. MILP Model | 8 |
| 3.2. Variables creation | 8 |
| 3.3. Constraints creation | 9 |
| 3.4. Optimal path | 11 |
| 3.5. Tests | 12 |
| 4. Part II - Simulated Annealing | 18 |
| 4.1. Implementation | 18 |
| 4.1.1. Initial solution | 18 |
| 4.1.2. Neighbour generation | 19 |
| 4.1.3. Cooling schedule | 19 |
| 4.1.4. Stopping criteria | 20 |
| 4.2. Parameters tuning | 20 |
| 4.2.1. TSP size: 10 | 21 |
| 4.2.2. TSP size: 25 | 21 |
| 4.2.3. TSP size: 50 | 22 |
| 4.2.4. TSP size: 75 | 22 |
| 4.2.5. TSP size: 100 | 23 |
| 4.3. Parameters Recap | 23 |
| 5. Performances comparison | 25 |
| 6. Conclusion | 27 |

Figures

| | |
|--|---|
| Figure 1: Project structure | 4 |
| Figure 2: Instance with 75 points | 6 |
| Figure 3: Instance with 100 points | 6 |
| Figure 4: Optimal path with 75 points | 7 |
| Figure 5: Optimal path with 100 points | 7 |

Tables

| | |
|---|----|
| Table 1: Cplex API Benchmarks | 12 |
| Table 2: Cplex API average performances | 17 |
| Table 3: SA performances for TSP size 10 | 21 |
| Table 4: SA performances for TSP size 25 | 21 |
| Table 5: SA performances for TSP size 50 | 22 |
| Table 6: SA performances for TSP size 75 | 22 |
| Table 7: SA performances for TSP size 100 | 23 |
| Table 8: SA parameters for different TSP size | 23 |
| Table 9: Performances comparison of SA and CPLEX Solver | 25 |

1. Introduction

This project presents two solutions for the Travelling Salesman Problem (TSP), utilizing two distinct approaches:

1. **Mathematical Model:** This approach employs the CPLEX C API's to formulate and solve the TSP as an optimization problem.
2. **Heuristic Model:** This approach uses the Simulated Annealing algorithm to find an approximate solution to the TSP.

Both methods aim to minimize the total travel distance while visiting each point exactly once and returning to the starting point. The mathematical model guarantees an optimal solution, whereas the heuristic model provides a near-optimal solution in a shorter amount of time.

1.1. Problem description

A company produces boards with holes used to build electric panels. Boards are positioned over a machine and a drill moves over the board, stops at the desired positions and makes the holes. Once a board is drilled, a new board is positioned and the process is iterated many times. Given the position of the holes on the board, the company asks us to determine the hole sequence that minimizes the total drilling time, taking into account that the time needed for making an hole is the same and constant for all the holes.

1.2. Project structure

The project is organized in the `TSP-solver` folder as shown in [Figure 1](#).

- `data/` : it contains the datasets generated by the script `data_generator.py`. In the subfolder `sol/` it stores all the path solutions for each dataset.
- `plots/` : it contains the plots generated by the script `data_generator.py`. In the subfolder `sol/` it stores all the plots with their path solution.
- `report.pdf` : the written report of the project.
- `scripts/` : it contains the python script used to generate and plot data. They are explained in section [2. Instances generator](#).
- `src/` : it contains the source files of the c++ project:
 - `main.cpp` : the main file of the project. It runs the CplexSolver and SASolver on two datasets.
 - `makefile` : the makefile used to compile the project.
 - `CplexSolver.cpp/.h` : a TSP solver which uses the CPLEX C API's. It implements the abstract class `TSPSolver`. It is explained in section [Part I: Cplex Solver](#).
 - `cpmacro.h` : useful macros of the CPLEX C API's.
 - `Point.cpp/.h` : a class representing a single point.
 - `SASolver.cpp/.h` : a TSP solver which uses the Simulated Annealing algorithm. It implements the abstract class `TSPSolver`. It is explained in section [Part II: Simulated Annealing](#).
 - `TSP.cpp/.h` : a TSP representation of the problem. It loads points from a dataset and calculates the matrix costs.

- `TSPSolution.cpp/.h` : a representation of a path solution to a TSP problem.
- `TSPSolver.cpp/.h` : an abstract class representing a solver for the TSP.

```

TSP-solver/
├── data/
│   ├── sol/
│   │   ├── Path_TSP_10_1.dat
│   │   └── ...
│   ├── TSP_10_1.dat
│   └── ...
├── plots/
│   ├── sol/
│   │   ├── Path_TSP_10_1.png
│   │   └── ...
│   ├── TSP_10_1.png
│   └── ...
├── report.pdf
├── scripts/
│   ├── data_generator.py
│   └── plot_solutions.py
└── src/
    ├── main.cpp
    ├── makefile
    ├── CplexSolver.cpp
    ├── CplexSolver.h
    ├── cpxmacro.h
    ├── point.cpp
    ├── point.h
    ├── SASolver.cpp
    ├── SASolver.h
    ├── TSP.cpp
    ├── TSP.h
    ├── TSPSolution.cpp
    ├── TSPSolution.h
    ├── TSPSolver.cpp
    └── TSPSolver.h

```

Figure 1: Project structure

1.3. Compilation

The project is written in C++ and compiles using the makefile provided in the `src` folder. To compile the project, use the command `make` in the `src/` folder. The makefile compiles the project using the `g++` compiler and produces an executable named `main`. The project uses the C++11 standard, so it is necessary to use a compiler that supports this standard.

It is possible to clean the project using the command `make clean`. This command removes the executable and all the object files.

The python scripts provided in the `scripts` folder and used to generate and plot data can be run using the command `python3 script_name.py`. The scripts require the `matplotlib` and `numpy` libraries to be installed. The libraries can be easily installed using the command `pip install matplotlib` and `pip install numpy`.

Note that any re run of the scripts will overwrite the existing data and plots. Since the data is generated randomly, the results may vary between runs.

All the tests have been made on the LabTA calculators from remote, on Ubuntu Operative System.

2. Instances generator

The instances were generated with the script `data_generator.py`. The script considers problems with 10, 25, 50, 75, and 100 points. For each problem size, it generates 9 examples. The number of examples is chosen considering the second part of the exercise: a significant portion of the training samples (66%) will be used to find the best parameters for the searches, while the remaining part will be used for testing.

To accurately recreate the problem description, the datasets are generated with specific criteria. Points are placed in shapes such as rectangles, squares, and circles. Shapes are picked randomly at each iteration. Additionally, there are points considered to be holes for bolts: there are at least 4 in each dataset, placed near the corners of the board.

Points are generated randomly according to a shape, but the script ensures that the distance between points is at least 0.3. This is done to avoid points being too close to each other, which could lead to numerical instability in the solution, particularly when subtracting numbers very close to each other.

The script also ensures that shapes do not overlap. Before adding a new shape, it checks for overlaps with existing shapes using specific functions:

- `check_overlap()` for circles
- `check_rectangle_overlap()` for rectangles and squares
- `check_circle_rectangle_overlap()` for checking overlaps between circles and rectangles

The data is stored in a `.dat` file, named `TSP_XX.X.dat`. The first line contains the number of nodes generated, and each subsequent line stores the coordinates of a point. The datasets are stored in the `/data` folder.

The script generates also plots to visualize data using the function `save_plot_as_file()`, which are stored in the `/plot` folder. Some examples can be seen in [Figure 2](#) and [Figure 3](#).

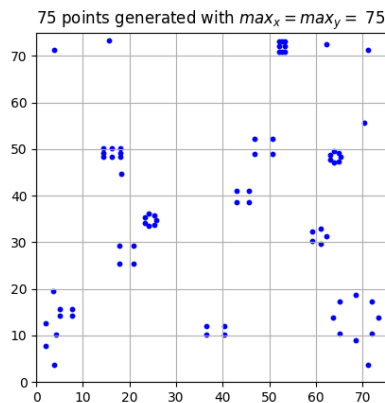


Figure 2: Instance with 75 points




Figure 3: Instance with 100 points

To visualize solutions, the script `plot_solutions.py` plots the dataset with the path solution. In particular, it reads data from a `.dat` file containing the dataset with the points coordinates placed in the `data/` folder and from another `.dat` file the path solution, placed in the `data/sol/` folder. Some examples can be visualized in [Figure 4](#) and [Figure 5](#).

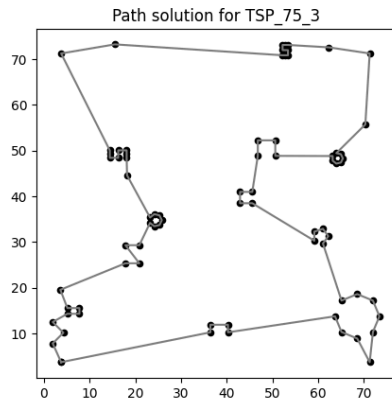


Figure 4: Optimal path with 75 points

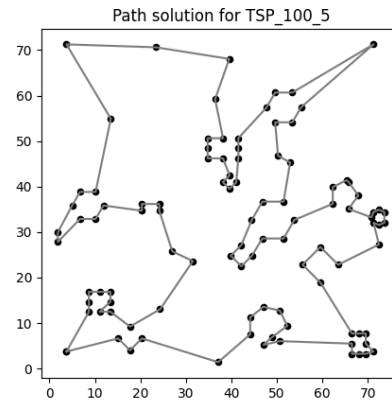


Figure 5: Optimal path with 100 points

3. Part I: Cplex Solver

This section details the implementation of the TSP model using the CPLEX API for C++, in the class `CplexSolver`.

3.1. MILP Model

Sets:

- N = graph nodes, representing the holes;
- A = arcs $(i, j), \forall i, j \in N$, representing the trajectory covered by the drill to move from hole i to hole j

Parameters:

- c_{ij} = time taken by the drill to move from i to $j, \forall (i, j) \in A$
- 0 = arbitrarily selected starting node, $0 \in N$.

Decision Variables:

- x_{ij} = amount of the flow shipped from i to $j, \forall (i, j) \in A$;
- $y_{ij} = 1$ if arc (i, j) ships some flow, 0 otherwise, $\forall (i, j) \in A$

Integer Linear Programming model:

$$\min \sum_{i,j:(i,j) \in A} c_{ij} y_{ij} \quad (1)$$

$$s.t. \sum_{i:(i,k) \in A} x_{ik} - \sum_{j:(k,j) \in A, j \neq 0} x_{kj} = 1 \quad \forall k \in N \setminus \{0\} \quad (2)$$

$$\sum_{j:(i,j) \in A} y_{ij} = 1 \quad \forall i \in N \quad (3)$$

$$\sum_{i:(i,j) \in A} y_{ij} = 1 \quad \forall j \in N \quad (4)$$

$$x_{ij} \leq (|N| - 1) y_{ij} \quad \forall (i, j) \in A, j \neq 0 \quad (5)$$

$$x_{ij} \in \mathbb{R}_+ \quad \forall (i, j) \in A, j \neq 0 \quad (6)$$

$$y_{ij} \in \{0, 1\} \quad \forall (i, j) \in A \quad (7)$$

3.2. Variables creation

Variables x and y are created using the `CPXnewcols` function with two calls, one for x variables and one for y variables.

To keep a connection between the CPLEX variables representation, indexes are stored in a map, defined as follows:

```
std::map<std::tuple<char, int, int>, int> vars_map;
```

The `CplexSolver` class implements the following private method to store the indexes in the map, where `type` is the variable type, `i` and `j` are the indexes of the variable and `cur_index` is the current index of the variable in the CPLEX model:

```
void add_index_to_map(char type, int i, int j, int cur_index)
```

To retrieve the CPLEX index of a variable, the following private method is implemented:

```
int get_index_from_map(char type, int i, int j)
```


To efficiently insert variables in the CPLEX model:

1. vectors are created to define: variable types, lower and upper bounds, objective coefficients and variable names. For each variable, the corresponding index is stored in the map using the `add_index_to_map()` method.
2. The `CPXnewcols` function is called once to add all y variables and then all x variables to the CPLEX problem. This is more efficient than adding variables one by one.
3. Memory cleanup is performed: the allocated space for variables names (pointers to string) is released after the variables are added to the CPLEX model.

3.3. Constraints creation

Similarly to the variables creation, the row constraints (2), (3), (4), (5) are added in a single call for each one of them: the `CPXaddrows` function is called once for each constraint type, adding all the constraints of that type to the CPLEX model.

To efficiently insert constraints in the CPLEX model.

1. vectors are initialized to define: right-hand side values, the sense of the constraints, the indexes of the variables involved, and their coefficients.
2. The `CPXaddrows` function is called once to add all the constraints at once. This is more efficient than adding constraints one by one.
3. Memory cleanup is performed: the allocated space for the constraints names (pointers to string) is released after the constraints are added to the CPLEX model.

Let's take a look at the values of the vectors used to add the constraints (2), (3), (4), (5) to the CPLEX model.

Constraint (2)

$$\sum_{i:(i,k) \in A} x_{ik} - \sum_{j:(k,j), j \neq 0} x_{kj} = 1 \quad \forall k \in N \setminus \{0\}$$

- `rowcnt = n - 1` is the number of constraints of this type. In fact we are considering all the nodes except the starting one.
- `nzcnt = rowcnt * (2*n - 3)` is the number of non-zero coefficients in the constraint matrix. In fact, for each node k we have $n - 1$ variables x_{ik} and $n - 2$ variables x_{kj} , because $j \neq 0$.
- `vector <double> rhs(rowcnt, 1)` is the right-hand side values of the constraints. In this case, all the values are set to 1.
- `vector <char> sense(rowcnt, 'E')` is the sense of the constraints. In this case, all the constraints are equalities, identified by char 'E'.
- `vector <int> rmatbeg(rowcnt)` is the starting index of the non-zero coefficients in the constraint matrix. In this case, the starting index of each row is given by $k * (2 * n - 3)$: in fact, at iteration k , we have already added $k * (n - 1)$ variables x_{ik} and $k * (n - 2)$ variables x_{kj} .
- `vector <int> rmatind(nzcnt)` is the indexes of the variables involved in the constraints. The indexes are stored in the vector at each iteration using the `get_index_from_map()` method. An internal counter is used to keep track of the current index in the vector.

- `vector <double> rmatval(nzcnt, 0)` is the coefficients of the variables involved in the constraints. The coefficients are set to 1 for the variables x_{ik} and -1 for the variables x_{kj} . As it happens for the indexes, the coefficients are stored in the vector using an internal counter to keep track of the current index.
- `vector <char*> rowname(rowcnt)` is the names of the constraints. The names are set using the `cpxString()` method which returns a pointer to a string after allocating the necessary space.

Constraint (3)

$$\sum_{j:(i,j) \in A} y_{ij} = 1 \quad \forall i \in N$$

- `rowcnt = n`, in fact we are creating a constraint for each node $i \in N$.
- `nzcnt = rowcnt * n`, in fact, for each row i we are summing n variables y .
- `vector <double> rhs(rowcnt, 1)`: in this case, all the right hand sides are set to 1.
- `vector <char> sense(rowcnt, 'E')` all the constraints are equalities, identified by char 'E'.
- `vector <int> rmatbeg(rowcnt)`: in this case, the starting index of each row i is given by $i * n$: in fact, at iteration i , we have already added $i * n$ variables y_{ij} .
- `vector <int> rmatind(nzcnt)` the indexes are stored in the vector at each iteration using `get_index_from_map()`, in particular the variable y_{ij} is in the position $i * n + j$ of this vector.
- `vector <double> rmatval(nzcnt, 1)` the coefficients are set to 1 for all the variables y_{ij} .
- `vector <char*> rowname(rowcnt)`: at each iteration i a new pointer to a string is added to this vector.

Constraint (4)

$$\sum_{i:(i,j) \in A} y_{ij} = 1 \quad \forall j \in N$$

- `rowcnt = n`, in fact we are creating a constraint for each node $j \in N$.
- `nzcnt = rowcnt * n`, in fact, for each row j we are summing n variables y .
- `vector <double> rhs(rowcnt, 1)`: in this case, all the right hand sides are set to 1.
- `vector <char> sense(rowcnt, 'E')` all the constraints are equalities, identified by char 'E'.
- `vector <int> rmatbeg(rowcnt)`: the starting index of each row j is given by $j * n$: in fact, at iteration j , we have already added $j * n$ variables y_{ij} .
- `vector <int> rmatind(nzcnt)` the indexes are stored in the vector at each iteration using `get_index_from_map()`, in particular the variable y_{ij} is in the position $j * n + i$ of this vector.
- `vector <double> rmatval(nzcnt, 1)` the coefficients are set to 1 for all the variables y_{ij} .

- `vector <char*> rowname(rowcnt)` : at each iteration j a new pointer to a string is added to this vector.

Constraint (5)

$$x_{ij} \leq (|N| - 1)y_{ij} \quad \forall (i, j) \in A, j \neq 0$$

Note that the constraint is added into the CPLEX model moving all the variables on the left-hand side:

$$x_{ij} \leq (|N| - 1)y_{ij} \quad \Rightarrow \quad x_{ij} - (|N| - 1)y_{ij} \leq 0$$

- `rowcnt = n * (n - 1)` , in fact we are creating a constraint for each pair of arc, except for the ones to the starting node.
- `nzcnt = 2 * rowcnt` , in fact, for each row ij we are summing two variables.
- `vector <double> rhs(rowcnt, 0)` : all the right hand sides are set to 0.
- `vector <char> sense(rowcnt, 'L')` all the constraints are less or equal, identified by char 'L'.
- `vector <int> rmatbeg(rowcnt)` : the starting index of each row constraint ij is given by $2 * row$, where row is an internal counter to keep track of the number of row constraints added.
- `vector <int> rmatind(nzcnt)` the indexes are stored in the vector at each iteration using `get_index_from_map()` , in particular, at a given row the variable x_{ij} is in position $row * 2$ and the variable y_{ij} is in position $row * 2 + 1$.
- `vector <double> rmatval(nzcnt)` the coefficients are set to 1 for all the variables x_{ij} , so the ones in position $row * 2$, and to $-(|N| - 1)$ for all the variables y_{ij} , so the ones in position $row * 2 + 1$.
- `vector <char*> rowname(rowcnt)` : at each iteration a new pointer to a string is added to this vector.

3.4. Optimal path

The optimal path is retrieved from the CPLEX model after the optimization process is completed. Using the CPLEX function `CPXgetx()` the values of the variables are stored in a private member of the `CplexSolver` class: `vector<double> varvals` .

Then, the optimal path is reconstructed by following the values of the variables y_{ij} using the private method `createPathSolution()` . It starts by setting the initial node as node 0 and adding it to the path. Then, it iteratively explores the subsequent nodes by checking, for each possible destination j , whether the variable y_{ij} indicates an active connection (when it is set to 1). When it finds a connected node j , it adds j to the path, updates the current node, and continues until all nodes have been visited. Once the tour is complete, it adds the starting node to the end of the path to represent the return to the starting point. In the end, the optimal path is stored in the `vector<int> path_solution` variable as an ordered sequence of nodes.

The optimal path is then returned by the public method `getPathSolution()` . This method is used to initialize a `TSPSolution` object, which is then used to print the path solution, like that:

```
>> Path solution: < 0 6 7 2 4 8 5 3 1 9 0 >
```

3.5. Tests

All the tests have been made in the LabTA calculators from remote, using Ubuntu operative system. The time limits were set to 1, 10, 30, 60, 120, 300 seconds. A cycle was made to collect the results of the tests and store them in a csv file. The results are shown in the following [Table 1](#) and [Table 2](#) figures.

Note that it was not possible to find an optimal path solution for the dataset `TSP_100_2.dat` within the time limit of 300 seconds.

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|--------------|--------------|------------|--------------|----------|-----------------|
| TSP_10_1.dat | 10 | 1 sec | 0.2577 sec | Optimal | 42.1044 |
| TSP_10_2.dat | 10 | 1 sec | 0.2497 sec | Optimal | 40.5850 |
| TSP_10_3.dat | 10 | 1 sec | 0.2306 sec | Optimal | 39.9835 |
| TSP_10_4.dat | 10 | 1 sec | 0.2850 sec | Optimal | 40.8729 |
| TSP_10_5.dat | 10 | 1 sec | 0.2702 sec | Optimal | 39.6283 |
| TSP_10_6.dat | 10 | 1 sec | 0.2279 sec | Optimal | 39.2479 |
| TSP_10_7.dat | 10 | 1 sec | 0.2481 sec | Optimal | 41.6571 |
| TSP_10_8.dat | 10 | 1 sec | 0.3347 sec | Optimal | 41.3277 |
| TSP_10_9.dat | 10 | 1 sec | 0.2816 sec | Optimal | 41.7131 |
| TSP_25_1.dat | 25 | 1 sec | 1.0065 sec | Feasible | 118.4207 |
| TSP_25_1.dat | 25 | 10 sec | 1.4000 sec | Optimal | 117.5848 |
| TSP_25_2.dat | 25 | 1 sec | 0.7831 sec | Optimal | 114.5453 |
| TSP_25_3.dat | 25 | 1 sec | 1.0128 sec | Feasible | 116.3888 |
| TSP_25_3.dat | 25 | 10 sec | 1.2669 sec | Optimal | 116.3888 |
| TSP_25_4.dat | 25 | 1 sec | 0.9641 sec | Optimal | 116.8552 |
| TSP_25_5.dat | 25 | 1 sec | 1.0213 sec | Feasible | 116.2516 |
| TSP_25_5.dat | 25 | 10 sec | 1.0392 sec | Optimal | 114.5004 |
| TSP_25_6.dat | 25 | 1 sec | 1.0104 sec | Feasible | 137.8108 |
| TSP_25_6.dat | 25 | 10 sec | 1.6192 sec | Optimal | 109.9704 |
| TSP_25_7.dat | 25 | 1 sec | 0.6487 sec | Optimal | 113.8732 |
| TSP_25_8.dat | 25 | 1 sec | 0.3907 sec | Optimal | 109.9128 |

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|--------------|--------------|------------|--------------|----------|-----------------|
| TSP_25_9.dat | 25 | 1 sec | 0.6006 sec | Optimal | 124.9924 |
| TSP_50_1.dat | 50 | 1 sec | 1.0174 sec | Feasible | 196.0226 |
| TSP_50_1.dat | 50 | 10 sec | 4.2766 sec | Optimal | 146.0800 |
| TSP_50_2.dat | 50 | 1 sec | 1.0107 sec | Feasible | 419.2273 |
| TSP_50_2.dat | 50 | 10 sec | 8.1265 sec | Optimal | 144.3220 |
| TSP_50_3.dat | 50 | 1 sec | 1.0087 sec | Feasible | 464.6008 |
| TSP_50_3.dat | 50 | 10 sec | 6.5681 sec | Optimal | 146.3909 |
| TSP_50_4.dat | 50 | 1 sec | 1.0131 sec | Feasible | 181.4801 |
| TSP_50_4.dat | 50 | 10 sec | 6.0332 sec | Optimal | 138.1332 |
| TSP_50_5.dat | 50 | 1 sec | 1.0103 sec | Feasible | 356.3201 |
| TSP_50_5.dat | 50 | 10 sec | 5.8066 sec | Optimal | 136.5960 |
| TSP_50_6.dat | 50 | 1 sec | 1.0116 sec | Feasible | 478.9975 |
| TSP_50_6.dat | 50 | 10 sec | 5.5610 sec | Optimal | 145.1683 |
| TSP_50_7.dat | 50 | 1 sec | 1.0133 sec | Feasible | 498.5362 |
| TSP_50_7.dat | 50 | 10 sec | 10.0246 sec | Feasible | 130.2013 |
| TSP_50_7.dat | 50 | 30 sec | 12.1113 sec | Optimal | 130.2013 |
| TSP_50_8.dat | 50 | 1 sec | 1.0123 sec | Feasible | 467.8946 |
| TSP_50_8.dat | 50 | 10 sec | 5.6500 sec | Optimal | 135.8707 |
| TSP_50_9.dat | 50 | 1 sec | 1.0136 sec | Feasible | 469.0935 |
| TSP_50_9.dat | 50 | 10 sec | 5.8418 sec | Optimal | 134.6858 |
| TSP_75_1.dat | 75 | 1 sec | 1.0097 sec | Feasible | 1459.4778 |
| TSP_75_1.dat | 75 | 10 sec | 10.0514 sec | Feasible | 507.3574 |
| TSP_75_1.dat | 75 | 30 sec | 30.0223 sec | Feasible | 450.7407 |
| TSP_75_1.dat | 75 | 60 sec | 37.2917 sec | Optimal | 450.5143 |
| TSP_75_2.dat | 75 | 1 sec | 1.0155 sec | Feasible | 1150.0906 |
| TSP_75_2.dat | 75 | 10 sec | 10.0691 sec | Feasible | 478.0992 |

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|--------------|--------------|------------|--------------|----------|-----------------|
| TSP_75_2.dat | 75 | 30 sec | 30.0389 sec | Feasible | 474.8855 |
| TSP_75_2.dat | 75 | 60 sec | 56.4847 sec | Optimal | 418.8488 |
| TSP_75_3.dat | 75 | 1 sec | 1.0196 sec | Feasible | 1206.3164 |
| TSP_75_3.dat | 75 | 10 sec | 10.1449 sec | Feasible | 571.3963 |
| TSP_75_3.dat | 75 | 30 sec | 30.0226 sec | Feasible | 534.4228 |
| TSP_75_3.dat | 75 | 60 sec | 51.6793 sec | Optimal | 389.0620 |
| TSP_75_4.dat | 75 | 1 sec | 1.0104 sec | Feasible | 1319.4935 |
| TSP_75_4.dat | 75 | 10 sec | 10.0252 sec | Feasible | 552.9256 |
| TSP_75_4.dat | 75 | 30 sec | 30.0244 sec | Feasible | 451.1586 |
| TSP_75_4.dat | 75 | 60 sec | 60.0304 sec | Feasible | 444.8608 |
| TSP_75_4.dat | 75 | 120 sec | 63.7516 sec | Optimal | 444.3585 |
| TSP_75_5.dat | 75 | 1 sec | 1.0107 sec | Feasible | 1158.4734 |
| TSP_75_5.dat | 75 | 10 sec | 10.0316 sec | Feasible | 533.9445 |
| TSP_75_5.dat | 75 | 30 sec | 30.0391 sec | Feasible | 458.2926 |
| TSP_75_5.dat | 75 | 60 sec | 34.0651 sec | Optimal | 458.2926 |
| TSP_75_6.dat | 75 | 1 sec | 1.0139 sec | Feasible | 1591.8847 |
| TSP_75_6.dat | 75 | 10 sec | 10.0198 sec | Feasible | 577.7397 |
| TSP_75_6.dat | 75 | 30 sec | 30.0305 sec | Feasible | 577.7397 |
| TSP_75_6.dat | 75 | 60 sec | 52.7740 sec | Optimal | 475.2658 |
| TSP_75_7.dat | 75 | 1 sec | 1.0251 sec | Feasible | 1492.1890 |
| TSP_75_7.dat | 75 | 10 sec | 10.0249 sec | Feasible | 623.1681 |
| TSP_75_7.dat | 75 | 30 sec | 30.0311 sec | Feasible | 417.7575 |
| TSP_75_7.dat | 75 | 60 sec | 40.6911 sec | Optimal | 415.6373 |
| TSP_75_8.dat | 75 | 1 sec | 1.0429 sec | Feasible | 1341.8736 |
| TSP_75_8.dat | 75 | 10 sec | 10.0205 sec | Feasible | 611.0727 |
| TSP_75_8.dat | 75 | 30 sec | 30.0282 sec | Feasible | 445.3625 |

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|---------------|--------------|------------|--------------|----------|-----------------|
| TSP_75_8.dat | 75 | 60 sec | 51.7496 sec | Optimal | 417.5275 |
| TSP_75_9.dat | 75 | 1 sec | 1.0298 sec | Feasible | 1332.3518 |
| TSP_75_9.dat | 75 | 10 sec | 10.0193 sec | Feasible | 581.9756 |
| TSP_75_9.dat | 75 | 30 sec | 30.0373 sec | Feasible | 449.0168 |
| TSP_75_9.dat | 75 | 60 sec | 43.3595 sec | Optimal | 449.0168 |
| TSP_100_1.dat | 100 | 1 sec | 1.0147 sec | Feasible | 1859.4852 |
| TSP_100_1.dat | 100 | 10 sec | 10.0569 sec | Feasible | 812.8012 |
| TSP_100_1.dat | 100 | 30 sec | 31.5660 sec | Feasible | 575.2419 |
| TSP_100_1.dat | 100 | 60 sec | 60.0495 sec | Feasible | 551.7561 |
| TSP_100_1.dat | 100 | 120 sec | 120.0306 sec | Feasible | 517.1964 |
| TSP_100_1.dat | 100 | 300 sec | 121.7011 sec | Optimal | 517.1964 |
| TSP_100_2.dat | 100 | 1 sec | 1.0151 sec | Feasible | 1579.5440 |
| TSP_100_2.dat | 100 | 10 sec | 10.0158 sec | Feasible | 836.5979 |
| TSP_100_2.dat | 100 | 30 sec | 30.1112 sec | Feasible | 836.5979 |
| TSP_100_2.dat | 100 | 60 sec | 60.2576 sec | Feasible | 504.9482 |
| TSP_100_2.dat | 100 | 120 sec | 120.0497 sec | Feasible | 479.1978 |
| TSP_100_2.dat | 100 | 300 sec | 300.0476 sec | Feasible | 473.2960 |
| TSP_100_3.dat | 100 | 1 sec | 1.0222 sec | Feasible | 1784.1474 |
| TSP_100_3.dat | 100 | 10 sec | 10.1491 sec | Feasible | 910.8624 |
| TSP_100_3.dat | 100 | 30 sec | 30.0316 sec | Feasible | 910.2184 |
| TSP_100_3.dat | 100 | 60 sec | 60.0432 sec | Feasible | 910.2184 |
| TSP_100_3.dat | 100 | 120 sec | 120.0332 sec | Feasible | 571.8121 |
| TSP_100_3.dat | 100 | 300 sec | 184.6020 sec | Optimal | 491.1374 |
| TSP_100_4.dat | 100 | 1 sec | 1.0118 sec | Feasible | 1592.8707 |
| TSP_100_4.dat | 100 | 10 sec | 10.0222 sec | Feasible | 658.7710 |
| TSP_100_4.dat | 100 | 30 sec | 30.5145 sec | Feasible | 609.9025 |

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|---------------|--------------|------------|--------------|----------|-----------------|
| TSP_100_4.dat | 100 | 60 sec | 60.0478 sec | Feasible | 506.5066 |
| TSP_100_4.dat | 100 | 120 sec | 120.0308 sec | Feasible | 506.1950 |
| TSP_100_4.dat | 100 | 300 sec | 140.4908 sec | Optimal | 480.4839 |
| TSP_100_5.dat | 100 | 1 sec | 1.0377 sec | Feasible | 1458.2235 |
| TSP_100_5.dat | 100 | 10 sec | 10.0210 sec | Feasible | 834.0529 |
| TSP_100_5.dat | 100 | 30 sec | 31.1143 sec | Feasible | 823.6652 |
| TSP_100_5.dat | 100 | 60 sec | 60.0387 sec | Feasible | 555.5668 |
| TSP_100_5.dat | 100 | 120 sec | 120.2293 sec | Feasible | 555.3637 |
| TSP_100_5.dat | 100 | 300 sec | 144.6668 sec | Optimal | 541.8146 |
| TSP_100_6.dat | 100 | 1 sec | 1.0195 sec | Feasible | 1534.8214 |
| TSP_100_6.dat | 100 | 10 sec | 10.4829 sec | Feasible | 687.1978 |
| TSP_100_6.dat | 100 | 30 sec | 30.0390 sec | Feasible | 686.8297 |
| TSP_100_6.dat | 100 | 60 sec | 60.0564 sec | Feasible | 603.9285 |
| TSP_100_6.dat | 100 | 120 sec | 120.3188 sec | Feasible | 536.9845 |
| TSP_100_6.dat | 100 | 300 sec | 191.0833 sec | Optimal | 477.2790 |
| TSP_100_7.dat | 100 | 1 sec | 1.0361 sec | Feasible | 1792.4031 |
| TSP_100_7.dat | 100 | 10 sec | 10.1240 sec | Feasible | 887.8424 |
| TSP_100_7.dat | 100 | 30 sec | 30.0509 sec | Feasible | 797.0106 |
| TSP_100_7.dat | 100 | 60 sec | 60.0497 sec | Feasible | 524.2325 |
| TSP_100_7.dat | 100 | 120 sec | 120.0249 sec | Feasible | 511.9271 |
| TSP_100_7.dat | 100 | 300 sec | 200.8718 sec | Optimal | 502.4607 |
| TSP_100_8.dat | 100 | 1 sec | 1.0200 sec | Feasible | 1870.0924 |
| TSP_100_8.dat | 100 | 10 sec | 10.0270 sec | Feasible | 992.9269 |
| TSP_100_8.dat | 100 | 30 sec | 30.0370 sec | Feasible | 992.9269 |
| TSP_100_8.dat | 100 | 60 sec | 60.0659 sec | Feasible | 577.8256 |
| TSP_100_8.dat | 100 | 120 sec | 88.9921 sec | Optimal | 478.1800 |

| Dataset | Problem size | Time limit | Solving Time | Status | Objective value |
|---------------|--------------|------------|--------------|----------|-----------------|
| TSP_100_9.dat | 100 | 1 sec | 1.0152 sec | Feasible | 1996.9046 |
| TSP_100_9.dat | 100 | 10 sec | 10.0219 sec | Feasible | 740.6538 |
| TSP_100_9.dat | 100 | 30 sec | 30.0269 sec | Feasible | 740.6538 |
| TSP_100_9.dat | 100 | 60 sec | 60.0634 sec | Feasible | 740.6538 |
| TSP_100_9.dat | 100 | 120 sec | 120.0594 sec | Feasible | 629.1568 |
| TSP_100_9.dat | 100 | 300 sec | 201.8372 sec | Optimal | 503.1572 |

Table 1: Cplex API Benchmarks

The average performances for the optimal path setup and solving time are shown in the following [Table 2](#).

| Problem size | Avg. time for setup | Avg. time for solving |
|--------------|---------------------|-----------------------|
| 10 nodes | 0.00074 sec | 0.265 sec |
| 25 nodes | 0.00181 sec | 0.963 sec |
| 50 nodes | 0.01511 sec | 6.857 sec |
| 75 nodes | 0.02768 sec | 48.101 sec |
| 100 nodes | 0.03779 sec | 141.551 sec |

Table 2: Cplex API average performances

4. Part II - Simulated Annealing

This section details the implementation of the TSP model using the Simulated Annealing algorithm, in the `SASolver` class. Simulated Annealing is a metaheuristic algorithm used to find the global optimum of a function. As the number of nodes in the TSP increases, the time required by the CPLEX solver to find the optimal solution also increases significantly. For this reason, a metaheuristic algorithm can provide a good solution within a reasonable time.

4.1. Implementation

The `SASolver` class is responsible for solving the TSP using the Simulated Annealing algorithm. In order to create a new instance of the solver, the user must provide the following parameters: the maximum number of iterations for each multistart, the cooling parameter, the initial temperature, the maximum number of non-improving iterations before stopping the algorithm. The algorithm is implemented in the `solve()` method, which takes as input the TSP instance and returns the best solution found.

The algorithm starts by generating an initial solution. Then, it iteratively generates new solutions by choosing a neighbor of the current solution and evaluating its cost. The new solution is accepted with a probability that depends on the difference in cost between the current and new solution and a temperature parameter that decreases over time. The algorithm follows a multirestart approach, meaning that it is executed multiple times with different initial solutions to improve the chances of finding a high-quality solution. It stops when the temperature reaches zero or a predefined number of iterations is reached. The best solution found across all runs is returned as the final solution.

4.1.1. Initial solution

Based on the number of iterations of multistart, different initial solutions are generated. The chosen number for the multistart is 5, so the algorithm will generate 5 different initial solutions. This approach is used to improve the chances of finding a high-quality solution. The value 5 represents a trade-off between small and large TSP instances. For smaller instances, too many multistarts would be unnecessary since the solution space is limited, and each SA run converges quickly. Conversely, for larger instances, multiple starting points help escape poor local optima. A proper tuning of the algorithm's parameters should ideally include optimizing the number of multistarts as well. However, due to time constraints, I have chosen to focus on tuning other parameters instead.

The initial solution, of the first iteration of the algorithm, is generated using the Nearest Neighbour heuristic. This heuristic uses a greedy approach: it starts from node 0 and iteratively selects the nearest unvisited node until all nodes are visited. The time complexity of this approach is $O(n^2)$. The method that is responsible for the Nearest Neighbour Initial solution is `initialNeighborSolution()`.

The second multistart of the algorithm uses the initial generated by `CheapestInsertionSolution()` method: starting from a tour with two nodes $\{0, 1, 0\}$, the algorithm iteratively selects the unvisited node that result in the minimum additional cost of inserting it in the tour. The time complexity of this implemented approach is $O(n^3)$ in the worst case.

The remaining multistarts of the algorithm use a random initial solution. The method that is responsible for the random initial solution is `randomize()`, of the `TSPSolution` class.

This approach allows the algorithm to explore different regions of the solution space: the Nearest Neighbour heuristic and the Cheapest Insertion heuristic provide a good starting point, while the random initial solution allows the algorithm to identify different local optima. In general, as said by the professor during the lectures, no evidences attest that better solutions are obtained starting from better initial solutions, however i chose to implement this approach in order to differentiate the origianl Simulated Annealing algorithm.

Note that the two heuristics were found online searching for TSP heuristics¹, however the implementations are made by myself.

4.1.2. Neighbour generation

The neighbourhood of a solution is defined as the set of solutions that can be obtained by applying the following moves:

- 2-opt move: reverse the order of the nodes between two randomly selected nodes. It is implemented in the `twoOptMove()` method; in order to evaluate the cost of the new solution, an override of the `evaluate()` method is used: given the two selected nodes i and j in the sequence $< 1...h, i, ..., j, l, ..., 1$, the method calculates the new costs of the solution using the following formula: $c_{\text{new}} = c_{\text{old}} - c_{hi} - c_{jl} + c_{hj} + c_{il}$
- 3-opt move: reverse the order of the nodes between three randomly selected nodes. It is implemented in the `threeOptMove()` method; in order to evaluate the cost of the new solution the standard evaluation method is used: starting from the initial node 0, the method calculates the cost of the new solution by summing the costs of the edges between the nodes in the sequence.
- swap move: swap the position of two randomly selected nodes. It is implemented in the `swapMove()` method; in order to evaluate the cost of the new solution the standard evaluation method is used, as for the 3-opt move.

At each iteration, one of the moves is selected at random and applied to the current solution to produce a neighboring solution. The method responsible for generating the new solution is `generateNeighbor()`. The moves have different probabilities of being selected: the distribution is based on the effectiveness of the 2-opt move in improving solutions, while the 3-opt and swap moves enhance exploration by escaping local optima. For large TSP instances, the 3-opt and swap moves are particularly useful, as they facilitate greater diversification in the solution space, while the 2-opt move is more effective for small instances. The probabilities of selecting the moves are defined in the `generateNeighbor()` method. Due to time constraints, I chose to set the probability of selecting the swap move to 0.4 and to fine tune the probabilities of selecting the 2-opt and 3-opt moves.

4.1.3. Cooling schedule

The temperature parameter is used to control the acceptance of new solutions. The temperature decreases over time according to a cooling parameter. The cooling schedule is th original one proposed by Kirkpatrick (1983): $T_{k+1} = \alpha * T_k$, where T_k is the temperature at iteration k and α is the cooling parameter. The temperature is initialized to a high value and decreases gradually to zero. The temperature is decreased at the end of each iteration.

¹[Some Important Heuristics for TSP, 2006](#)

4.1.4. Stopping criteria

The algorithm stops when a predefined number of iterations is reached, depending on the problem size.

4.2. Parameters tuning

The Simulated Annealing algorithm has several parameters that can be tuned to improve its performance.

As said before in the section [2. Instances generator](#), the training set is composed of 66% of the generated instances, while the remaining 33% is used for testing. The training set is used to find the best parameters for the Simulated Annealing algorithm. For each problem size, instances from 1 to 6 are used for training, while instances from 7 to 9 are used for testing. The algorithm is executed five times for each instance and the results are reported in different tables.

In general, the goal was to find the best parameters for the Simulated Annealing algorithm that would allow it to find a solution closer to the optimal one in a reasonable amount of time.

The parameters that are tuned are:

- N_{max} : the maximum number of iterations for each multistart
- α : the cooling parameter
- T : the initial temperature
- N_{bad} : the maximum number of non-improving iterations before stopping the algorithm
- P_{2opt} : the probability of selecting the 2-opt move at each iteration
- P_{3opt} : the probability of selecting the 3-opt move at each iteration

Due to the exponential growth in the number of tests required, an exhaustive parameter search was not feasible. Instead, a set of empirical tests was performed, and the best-performing combination of parameters is reported along with the corresponding results. In general the following considerations were made:

- N_{max} : the maximum number of iterations for each multistart should be proportional to the problem size. A higher number of iterations would allow the algorithm to explore the solution space more freely, while a lower number would make the algorithm act like a local search. As the number of iterations increases, the computational time also increases.
- α : the cooling parameter influences the convergence speed of the algorithm. For smaller instances, a higher value of alpha would allow the algorithm to converge faster, while for larger instances, a lower value would be more appropriate to allow the algorithm to explore the solution space more freely.
- T : the initial temperature should be proportional to the problem size. A lower temperature would make the algorithm act like a local search, accepting only improving solutions, while a higher temperature would allow the algorithm to explore the solution space more freely but with a higher probability of accepting worsening solutions.
- N_{bad} : the maximum number of non-improving iterations before stopping the algorithm should be proportional to the problem size. For bigger instances, an higher number of

non-improving iteration would allow the algorithm to find better solutions, because the solution space is larger and the probabilities of improving solutions are lower.

- P_{2opt} : is the the best move for small instances, because it is the most effective in improving solutions, because it is the most effective in improving solutions.
- P_{3opt} : is the best move for larger instances, because it allows the algorithm to escape local optima and explore the solution space more freely.

4.2.1. TSP size: 10

| N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|-----------|----------|-----|-----------|------------|------------|
| 500 | 0.99 | 12 | 200 | 0.2 | 0.4 |

| Training | | | | | |
|--------------|-----------|------------|---------|-----------|------------|
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_10_1.dat | 42.1044 | 42.1044 | 42.1044 | 0.0000 | 0.1421 sec |
| TSP_10_2.dat | 40.5850 | 40.9087 | 40.6498 | 0.1295 | 0.1611 sec |
| TSP_10_3.dat | 39.9835 | 39.9835 | 39.9835 | 0.0000 | 0.1464 sec |
| TSP_10_4.dat | 40.8729 | 41.1680 | 40.9319 | 0.1180 | 0.1157 sec |
| TSP_10_5.dat | 39.6283 | 39.6283 | 39.6283 | 0.0000 | 0.1793 sec |
| TSP_10_6.dat | 39.2479 | 39.2479 | 39.2479 | 0.0000 | 0.1595 sec |
| Testing | | | | | |
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_10_7.dat | 41.6571 | 41.6571 | 41.6571 | 0.0000 | 0.1521 sec |
| TSP_10_8.dat | 41.3277 | 41.3277 | 41.3277 | 0.0000 | 0.1334 sec |
| TSP_10_9.dat | 41.7131 | 42.0401 | 41.8439 | 0.1602 | 0.1627 sec |

Table 3: SA performances for TSP size 10

4.2.2. TSP size: 25

| N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|-----------|----------|-----|-----------|------------|------------|
| 1000 | 0.96 | 50 | 4000 | 0.15 | 0.45 |

| Training | | | | | |
|--------------|-----------|------------|----------|-----------|------------|
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_25_1.dat | 121.0208 | 127.7170 | 124.6791 | 2.1436 | 0.4801 sec |
| TSP_25_2.dat | 116.9371 | 120.4846 | 118.5675 | 1.2130 | 0.5295 sec |
| TSP_25_3.dat | 119.5117 | 124.4823 | 122.4935 | 2.4134 | 0.4263 sec |
| TSP_25_4.dat | 118.3463 | 124.6497 | 120.9251 | 2.1064 | 0.5039 sec |
| TSP_25_5.dat | 118.2204 | 120.7781 | 119.4260 | 0.8774 | 0.5577 sec |

| TSP_25_6.dat | 112.7662 | 114.5473 | 113.9544 | 0.6400 | 0.5026 sec |
|--------------|-----------|------------|----------|-----------|------------|
| Testing | | | | | |
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_25_7.dat | 114.5640 | 125.2926 | 118.6536 | 3.8626 | 0.4637 sec |
| TSP_25_8.dat | 109.9128 | 120.6628 | 114.4376 | 3.7549 | 0.5222 sec |
| TSP_25_9.dat | 131.0302 | 136.9749 | 134.3330 | 2.1988 | 0.5286 sec |

Table 4: SA performances for TSP size 25

4.2.3. TSP size: 50

| N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|-----------|----------|-----|-----------|------------|------------|
| 3000 | 0.94 | 100 | 1000 | 0.15 | 0.45 |

| Training | | | | | |
|--------------|-----------|------------|----------|-----------|------------|
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_50_1.dat | 158.2897 | 167.0534 | 161.9167 | 2.9616 | 3.0452 sec |
| TSP_50_2.dat | 158.7165 | 163.5184 | 161.1484 | 2.0541 | 3.3085 sec |
| TSP_50_3.dat | 156.1423 | 164.1950 | 160.3999 | 2.8786 | 3.0735 sec |
| TSP_50_4.dat | 149.3260 | 149.3260 | 149.3260 | 0.0000 | 2.8179 sec |
| TSP_50_5.dat | 142.0716 | 159.0605 | 148.5610 | 5.9536 | 3.2037 sec |
| TSP_50_6.dat | 152.0125 | 153.9758 | 153.5831 | 0.7853 | 3.2483 sec |
| Testing | | | | | |
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_50_7.dat | 135.9899 | 143.3034 | 140.0773 | 2.8414 | 3.2215 sec |
| TSP_50_8.dat | 146.4733 | 159.4977 | 152.7684 | 5.2280 | 3.1128 sec |
| TSP_50_9.dat | 144.1011 | 156.2556 | 151.2625 | 4.4977 | 3.2332 sec |

Table 5: SA performances for TSP size 50

4.2.4. TSP size: 75

| N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|-----------|----------|-----|-----------|------------|------------|
| 15000 | 0.92 | 250 | 5000 | 0.1 | 0.5 |

| Training | | | | | |
|--------------|-----------|------------|----------|-----------|-------------|
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_75_1.dat | 481.8080 | 496.1988 | 489.0034 | 7.1954 | 16.0154 sec |
| TSP_75_2.dat | 442.5938 | 453.9232 | 448.2585 | 5.6647 | 15.5979 sec |
| TSP_75_3.dat | 402.5947 | 448.8334 | 425.7140 | 23.1194 | 17.7486 sec |

| TSP_75_4.dat | 460.0896 | 477.1074 | 468.5985 | 8.5089 | 17.4065 sec |
|--------------|-----------|------------|----------|-----------|-------------|
| TSP_75_5.dat | 495.8942 | 501.2556 | 498.5749 | 2.6807 | 16.0858 sec |
| TSP_75_6.dat | 513.6851 | 556.0474 | 534.8663 | 21.1811 | 18.8757 sec |
| Testing | | | | | |
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_75_7.dat | 451.2295 | 455.2584 | 453.2440 | 2.0145 | 15.5708 sec |
| TSP_75_8.dat | 467.2923 | 477.4068 | 472.3495 | 5.0573 | 16.8591 sec |
| TSP_75_9.dat | 481.4110 | 503.1316 | 492.2713 | 10.8603 | 15.8453 sec |

Table 6: SA performances for TSP size 75

4.2.5. TSP size: 100

| N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|-----------|----------|-----|-----------|------------|------------|
| 25000 | 0.91 | 350 | 10000 | 0.1 | 0.5 |

| Training | | | | | |
|---------------|-----------|------------|----------|-----------|-----------|
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_100_1.dat | 562.0855 | 616.8389 | 589.4622 | 27.3767 | 32.8537 |
| TSP_100_2.dat | 517.6495 | 523.4737 | 520.5616 | 2.9121 | 36.3481 |
| TSP_100_3.dat | 538.9852 | 554.0208 | 546.5030 | 7.5178 | 38.6829 |
| TSP_100_4.dat | 511.0467 | 534.6725 | 522.8596 | 11.8129 | 36.0500 |
| TSP_100_5.dat | 604.2694 | 612.9694 | 608.6194 | 4.3500 | 33.7766 |
| TSP_100_6.dat | 524.8286 | 561.9574 | 543.3930 | 18.5644 | 33.5939 |
| Testing | | | | | |
| Instance | Best Sol. | Worst Sol. | Avg. | Std. Dev. | Avg. Time |
| TSP_100_7.dat | 557.1537 | 557.1537 | 557.1537 | 0.0000 | 32.9219 |
| TSP_100_8.dat | 518.6339 | 585.9742 | 552.3040 | 33.6702 | 32.3897 |
| TSP_100_9.dat | 568.9786 | 610.3024 | 589.6405 | 20.6619 | 31.9742 |

Table 7: SA performances for TSP size 100

4.3. Parameters Recap

| Size | N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|------|-----------|----------|-----|-----------|------------|------------|
| 10 | 500 | 0.99 | 12 | 200 | 0.2 | 0.4 |
| 25 | 1000 | 0.96 | 50 | 4000 | 0.15 | 0.45 |
| 50 | 3000 | 0.94 | 100 | 1000 | 0.15 | 0.45 |
| 75 | 15000 | 0.92 | 250 | 5000 | 0.1 | 0.5 |

| Size | N_{max} | α | T | N_{bad} | P_{2opt} | P_{3opt} |
|------|-----------|----------|-----|-----------|------------|------------|
| 100 | 25000 | 0.91 | 350 | 10000 | 0.1 | 0.5 |

Table 8: SA parameters for different TSP size

5. Performances comparison

[Table 9](#) compares the performances of the Cplex solver and the SA algorithm. The table shows that the Cplex implementation is always able to find the optimal solution, while the SA algorithm is not. However, the SA algorithm is able to find a resonable solution in small amount of time, even for the largest instances. The table also shows that the SA algorithm is able to find a solution that is close to the optimal one, with a relative error of less than $\approx 10\%$ for all instances.

| Instance | Cplex Sol. | Cplex Solving Time | SA Sol. | SA Solving Time | Gap % | Speedup SA |
|--------------|------------|--------------------|----------|-----------------|---------|------------|
| TSP_10_1.dat | 42.1044 | 0.2577 | 42.1044 | 0.1421 | 0.0 | 1.81 |
| TSP_10_2.dat | 40.585 | 0.2497 | 40.585 | 0.1611 | 0.0 | 1.55 |
| TSP_10_3.dat | 39.9835 | 0.2306 | 39.9835 | 0.1464 | 0.0 | 1.58 |
| TSP_10_4.dat | 40.8729 | 0.285 | 40.8729 | 0.1157 | 0.0 | 2.46 |
| TSP_10_5.dat | 39.6283 | 0.2702 | 39.6283 | 0.1793 | 0.0 | 1.51 |
| TSP_10_6.dat | 39.2479 | 0.2279 | 39.2479 | 0.1595 | 0.0 | 1.43 |
| TSP_10_7.dat | 41.6571 | 0.2481 | 41.6571 | 0.1521 | 0.0 | 1.63 |
| TSP_10_8.dat | 41.3277 | 0.3347 | 41.3277 | 0.1334 | 0.0 | 2.51 |
| TSP_25_1.dat | 117.5848 | 1.4 | 121.0208 | 0.4801 | -2.8392 | 2.92 |
| TSP_25_2.dat | 114.5453 | 0.7831 | 116.9371 | 0.5295 | -2.0454 | 1.48 |
| TSP_25_3.dat | 116.3888 | 1.2669 | 119.5117 | 0.4263 | -2.613 | 2.97 |
| TSP_25_4.dat | 116.8552 | 0.9641 | 118.3463 | 0.5039 | -1.2599 | 1.91 |
| TSP_25_5.dat | 114.5004 | 1.0392 | 118.2204 | 0.5577 | -3.1467 | 1.86 |
| TSP_25_6.dat | 109.9704 | 1.6192 | 112.7662 | 0.5026 | -2.4793 | 3.22 |
| TSP_25_7.dat | 113.8732 | 0.6487 | 114.564 | 0.4637 | -0.603 | 1.4 |
| TSP_25_8.dat | 109.9128 | 0.3907 | 109.9128 | 0.5222 | 0.0 | 0.75 |
| TSP_25_9.dat | 124.9924 | 0.6006 | 131.0302 | 0.5286 | -4.6079 | 1.14 |
| TSP_50_1.dat | 146.08 | 4.2766 | 158.2897 | 3.0452 | -7.7135 | 1.4 |
| TSP_50_2.dat | 144.322 | 8.1265 | 158.7165 | 3.3085 | -9.0693 | 2.46 |
| TSP_50_3.dat | 146.3909 | 6.5681 | 156.1423 | 3.0735 | -6.2452 | 2.14 |
| TSP_50_4.dat | 138.1332 | 6.0332 | 149.326 | 2.8179 | -7.4955 | 2.14 |
| TSP_50_5.dat | 136.596 | 5.8066 | 142.0716 | 3.2037 | -3.8541 | 1.81 |
| TSP_50_6.dat | 145.1683 | 5.561 | 152.0125 | 3.2483 | -4.5024 | 1.71 |
| TSP_50_7.dat | 130.2013 | 12.1113 | 135.9899 | 3.2215 | -4.2566 | 3.76 |
| TSP_50_8.dat | 135.8707 | 5.65 | 146.4733 | 3.1128 | -7.2386 | 1.82 |

| Instance | Cplex Sol. | Cplex Solving Time | SA Sol. | SA Solving Time | Gap % | Speedup SA |
|---------------|------------|--------------------|----------|-----------------|----------|------------|
| TSP_50_9.dat | 134.6858 | 5.8418 | 144.1011 | 3.2332 | -6.5338 | 1.81 |
| TSP_75_1.dat | 450.5143 | 37.2917 | 481.808 | 16.0154 | -6.4951 | 2.33 |
| TSP_75_2.dat | 418.8488 | 56.4847 | 442.5938 | 15.5979 | -5.365 | 3.62 |
| TSP_75_3.dat | 389.062 | 51.6793 | 402.5947 | 17.7486 | -3.3614 | 2.91 |
| TSP_75_4.dat | 444.3585 | 63.7516 | 460.0896 | 17.4065 | -3.4191 | 3.66 |
| TSP_75_5.dat | 458.2926 | 34.0651 | 495.8942 | 16.0858 | -7.5826 | 2.12 |
| TSP_75_6.dat | 475.2658 | 52.774 | 513.6851 | 18.8757 | -7.4792 | 2.8 |
| TSP_75_7.dat | 415.6373 | 40.6911 | 451.2295 | 15.5708 | -7.8878 | 2.61 |
| TSP_75_8.dat | 417.5275 | 51.7496 | 467.2923 | 16.8591 | -10.6496 | 3.07 |
| TSP_75_9.dat | 449.0168 | 43.3595 | 481.411 | 15.8453 | -6.729 | 2.74 |
| TSP_100_1.dat | 517.1964 | 121.7011 | 562.0855 | 32.8537 | -7.9862 | 3.7 |
| TSP_100_2.dat | 473.296 | 300.0476 | 517.6495 | 36.3481 | -8.5682 | 8.25 |
| TSP_100_3.dat | 491.1374 | 184.602 | 538.9852 | 38.6829 | -8.8774 | 4.77 |
| TSP_100_4.dat | 480.4839 | 140.4908 | 511.0467 | 36.05 | -5.9804 | 3.9 |
| TSP_100_5.dat | 541.8146 | 144.6668 | 604.2694 | 33.7766 | -10.3356 | 4.28 |
| TSP_100_6.dat | 477.279 | 191.0833 | 524.8286 | 33.5939 | -9.06 | 5.69 |
| TSP_100_7.dat | 502.4607 | 200.8718 | 557.1537 | 32.9219 | -9.8165 | 6.1 |
| TSP_100_8.dat | 478.18 | 88.9921 | 518.6339 | 32.3897 | -7.8001 | 2.75 |
| TSP_100_9.dat | 503.1572 | 201.8372 | 568.9786 | 31.9742 | -11.5683 | 6.31 |

Table 9: Performances comparison of SA and Cplex Solver

6. Conclusion

This exercise showed that the TSP is a complex problem that becomes increasingly challenging to solve as its size grows. The CPLEX solver can find the optimal solution for all instances, but it fails to do so within a reasonable amount of time for the largest instances. On the other hand, the Simulated Annealing algorithm can quickly provide high-quality solutions, even for the largest problem sizes. The implemented SA consistently finds solutions within $\approx 10\%$ of the optimal value, demonstrating that it is a viable alternative to CPLEX for large-scale TSP instances.

The effectiveness of SA heavily depends on parameter tuning, as the results are highly sensitive to the chosen values. A trade-off must be established between solution quality and computational time: in some cases, prioritizing a faster execution may be more important, while in others, achieving the best possible solution within a reasonable timeframe may be preferred. In general, the parameters can be adjusted following the suggestions provided in the section [4.2. Parameter tuning](#).